Toronto Housing Market Analysis using Business Intelligence

Executive Summary

This project addresses a critical challenge in Toronto's CAD $90 billion annual real estate market: the lack of transparent, accessible, data-driven insights for property pricing and valuation. Despite the wealth of listing data maintained by RE/MAX and other major brokerages through the Multiple Listing Service (MLS), buyers, sellers, and real estate professionals struggle with fragmented information sources, pricing opacity, and the inability to integrate macroeconomic factors like interest rates into property assessments. In a market where home prices have surged 156% over the past decade, this information asymmetry creates significant financial risk and market inefficiency.

The Problem: Toronto housing prices vary dramatically based on location, property type, number of bedrooms and bathrooms, and square footage, yet this information remains difficult to interpret when dealing with thousands of listings across diverse neighborhoods. Furthermore, external economic factors like Bank of Canada interest rates which directly influence mortgage affordability and buyer demand are not typically visualized alongside listing data, preventing stakeholders from understanding the complete market picture.

The BI Solution: This project implements a comprehensive Business Intelligence framework that transforms raw real estate data into actionable insights through a systematic five-stage pipeline.

Key Findings: The analysis revealed several critical market dynamics. Location analysis identified the top 10 most expensive FSA zones with statistically reliable sample sizes, providing buyers with geographic targets for value optimization. Market composition analysis showed that condominiums represent the highest volume of listings, but single-family homes command the highest median prices, helping buyers and investors understand inventory distribution.

Challenges and Limitations: Implementation faced several obstacles including data quality issues with corrupt property tax entries requiring extensive filtering, geographic data inconsistency necessitating the use of FSA codes rather than neighborhood names, and the inability to directly integrate interest rate data with listings due to missing consistent listing dates. The static nature of the dataset also means insights may become outdated as market

conditions evolve, highlighting the need for automated data refresh mechanisms in production deployments.

Final Outcome: This project successfully demonstrates that Business Intelligence methodologies can transform the Toronto real estate market from an opaque, experience-driven ecosystem into a transparent, data-driven marketplace accessible to all participants regardless of technical expertise.

Table of Contents

Executive Summary 2

1. Introduction 4

- Company Overview 4
- What is Business Intelligence? 4
- Why is BI Critical for Real Estate? 5
- Project Purpose 5

2. Problem Description 6

3. BI Solution / System Implementation 6

- 3.1 Data Sources (Data Acquisition) 6
- 3.2 Data Cleaning and Storage 7
- 3.3 Data Analysis and Visualization 7
  - Location Analysis 8
  - Market Composition 8
  - Property Configuration 9
  - Feature Drivers 9
- 3.4 Dashboard Implementation 10

4. Benefits 10

5. Challenges 11

6. Conclusion and Recommendations 12

7. References 14

Appendix: Work Breakdown Structure (WBS) 13

1. Introduction

Company Overview

The Toronto real estate market is Canada's largest property ecosystem, generating over 100,000 transactions annually worth approximately CAD $90 billion across the Greater Toronto Area's 6.4 million residents. This market encompasses a diverse range of property types from high-rise condominiums in the downtown core to detached single-family homes in suburban neighborhoods like North York, Scarborough, and Etobicoke.

RE/MAX, one of North America's leading real estate networks with over 140,000 agents globally, maintains one of the most comprehensive Multiple Listing Service (MLS) databases in Canada, capturing granular property-level data including street addresses, geographic coordinates, listing prices, bedroom and bathroom counts, square footage, property types, and listing dates. However, despite the wealth of available data, the Toronto housing market faces persistent challenges and inefficiencies that hinder effective decision-making for all stakeholders. Buyers must manually cross-reference fragmented listings without integrated comparison tools. Sellers struggle to price properties accurately due to limited neighborhood trend data. Historical price movements and seasonal patterns remain inaccessible to non-professionals. Most critically, external economic factors like Bank of Canada interest rates which directly impact mortgage affordability are disconnected from listing data, making it difficult to assess true market conditions. These challenges create information asymmetry that disadvantages first-time buyers, individual sellers, and small agencies while presenting a compelling case for Business Intelligence transformation.

WHAT IS BUSINESS INTELLIGENCE?

Raw Data	BI System	Actionable Insights

Business Intelligence (BI) is an integrated set of technologies and processes that transform raw data into actionable insights through data acquisition, cleaning, storage, analysis, and visualization. BI systems enable organizations to move from anecdotal decision-making to evidence-based strategy by revealing patterns, predicting trends, and quantifying outcomes. In real estate, BI turns thousands of fragmented listings into clear pricing benchmarks, neighborhood comparisons, and market forecasts of insights previously accessible only to institutional investors with dedicated analytics teams.

WHY IS BI CRITICAL FOR REAL ESTATE

Toronto home prices have surged 156% over the past decade, it is essential for market participants to navigate volatility, assess true value, and mitigate financial risk with comprehensive data analysis.
For **homebuyers**, a single property purchase represents the largest financial commitment most Canadians will ever make, yet information asymmetry forces them to rely on limited comparable sales, realtor opinions, or gut instinct rather than comprehensive market analysis, potentially resulting in overpayment by tens of thousands of dollars or missed opportunities in undervalued neighborhoods. For **sellers**, pricing a property without data-driven insights leads to either undervaluation that leaves money on the table or overpricing that results in extended time on market, price reductions, and stigmatization. For **real estate agents**, manually analyzing thousands of listings to provide clients with competitive market analyses is time-intensive, error-prone, and unsustainable at scale, while lacking the sophistication to integrate macroeconomic indicators like interest rate trends into valuation models. For **investors and developers**, identifying emerging neighborhoods, forecasting price appreciation, and optimizing portfolio allocation requires statistical rigor that manual analysis cannot provide.

Business Intelligence democratizes market intelligence, enabling transparent, efficient, and equitable participation across all stakeholders in Toronto's $90 billion housing market.

PROJECT PURPOSE

This project implements a comprehensive Business Intelligence solution to address the critical problem of pricing opacity and information fragmentation in the Toronto real estate market.
The approach follows five systematic steps:

1. Data Acquisition: Integrating Kaggle housing listings with Bank of Canada interest rate data via API.
2. Data Processing: Removing duplicates, handling missing values, filtering extreme outliers (e.g., $888 million property tax errors)
3. Data Storage: Organizing cleaned data in a SQLite database
4. Analytical Modeling: Conducting location analysis (using Forward Sortation Area postal codes), market composition analysis (property types and price distribution), property configuration analysis (bedroom/bathroom impact on price), and feature correlation analysis
5. Visualization: Delivering insights through an interactive five-page Power BI dashboard with dynamic filters for property type, price range, and neighborhood.

The goal is transforming inaccessible real estate data into actionable intelligence that empowers buyers to identify fair value, enables sellers to price competitively, and supports agents with data-backed recommendations ultimately demonstrating how BI can create more transparent and efficient housing markets.

2. Problem Description

The main problem addressed in this project is the lack of clear understanding of housing price patterns in the Toronto real estate market. Property prices vary significantly based on factors such as location, number of bedrooms and bathrooms, property type, and size. However, this information is not always easy to interpret, especially when dealing with a large number of listings.

One of the key challenges is that similar properties can have very different prices. For example, two houses with the same number of bedrooms may be priced differently due to differences in location or property type. Without proper analysis, it becomes difficult for buyers and real estate agents to determine whether a property is fairly priced.

In addition, the housing market generates a large volume of data, which makes manual analysis inefficient and unreliable. Realtors often rely on personal experience or limited comparisons instead of using complete data-driven insights. This can lead to poor decisions, such as overpricing or underpricing properties.

External economic factors, such as interest rates, also influence housing demand and pricing trends. However, these factors are not always clearly connected to listing data, making it harder to understand the overall market situation.

This problem is important because it affects buyers, sellers, and real estate agents. Buyers want to find good value, sellers want to price competitively, and agents need accurate insights to guide their clients.

Therefore, there is a need for a Business Intelligence solution that can analyze Toronto housing data, identify pricing patterns, and provide data driven insights to support better decision making.

3. BI Solution / System Implementation

3.1 Data Sources (Data Acquisition)

The primary dataset for this project was obtained from a publicly available housing listings dataset on Kaggle, which contains property-level information across Ontario. This dataset includes key variables such as address, locality, postal code, latitude and longitude, listing price, number of bedrooms and bathrooms, square footage, and property type. The dataset was uploaded into Google Colab and used as the foundation for the analysis, as it provides detailed insights into individual housing listings. These variables enable the examination of pricing

patterns, property characteristics, and regional differences, which are essential for supporting real estate decision-making.

To enhance the analysis with external market context, additional data was acquired using the Bank of Canada's Valet API. This API was used to retrieve official interest rate data, specifically the overnight rate, which plays a significant role in influencing mortgage rates and housing demand. The API data was accessed programmatically using Python's requests library and converted into a structured format for further use. While this dataset is not directly joined with the housing listings due to the absence of consistent listing dates, it provides valuable macroeconomic context that can support interpretation of pricing patterns and market conditions within the BI solution.

3.2 Data Cleaning and Storage

Before analysis, the collected housing dataset was cleaned to improve data quality and ensure reliable results. During the cleaning process, duplicate records were removed to prevent repeated entries from affecting the analysis. Missing values were handled by removing records with missing essential fields such as price, while other missing values, such as square footage, were filled using median values to maintain consistency. In addition, column names were standardized and important fields such as price, number of bedrooms, bathrooms, and square footage were converted into numeric formats to support accurate calculations.

After the data cleaning process, the dataset was stored in a structured format for easy access and future analysis. In this project, the cleaned data was stored in a **SQLite database**, which allows for efficient storage and retrieval of housing information. Storing the cleaned dataset in a database ensures that the data is organized, secure, and ready to be used for reporting and visualization in later stages of the Business Intelligence process.

3.3 Data Analysis and Visualization

The analysis of Toronto housing dataset follows a descriptive and diagnostic approach. Descriptive analysis was used to summarize overall market composition, pricing patterns, and inventory configurations. Diagnostic analysis was used to analyze the relationships between property features and listing prices. Median price was used as the primary measure throughout all visualizations to reduce the influence of extreme outliers, as the dataset contains a small number of luxury properties that would significantly distort a mean-based analysis. All visualizations were developed in Python using Matplotlib and Seaborn, and the results were subsequently presented through an interactive Power BI dashboard.

The analysis was structured into four areas:

Location Analysis

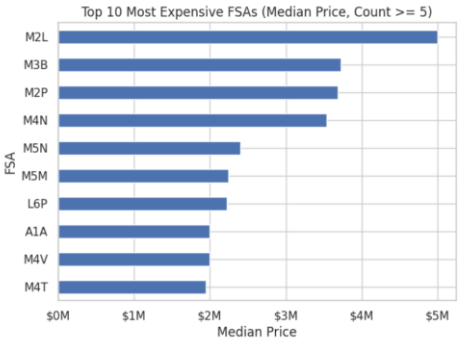

For location analysis, we examined how listing prices vary across Toronto neighborhoods. Although the dataset includes an address locality column, FSA (the first three characters of the postal code) was used instead, as it provides more consistent geographic groupings and is a standard approach in Canadian real estate analysis.

This allowed us to compare median prices across different areas of Toronto in a reliable way. A horizontal bar chart was created to rank the top 10 most expensive FSAs by median listing price, with a minimum threshold of five listings per FSA applied to ensure statistical reliability. A geographic scatter map was also produced to visualize the spatial distribution of listings and price levels across the city.

Market Composition

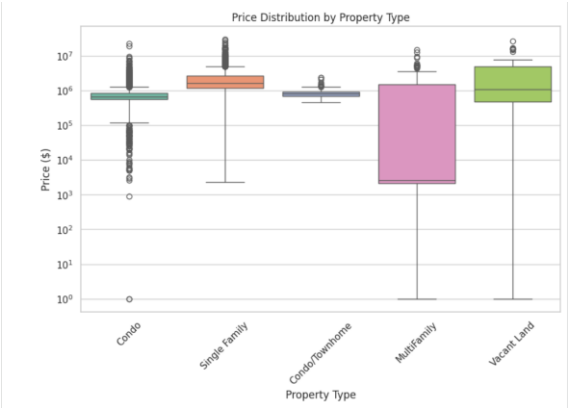

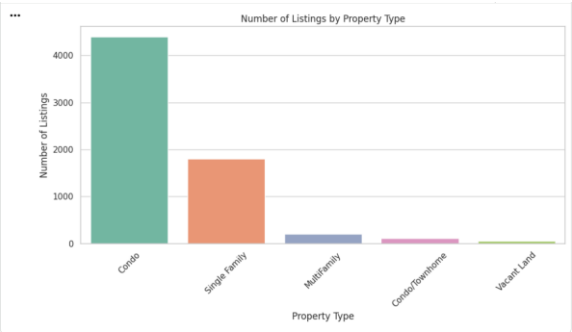

The second area, market composition, looked at price distribution and property type patterns. A histogram was created to show how listing prices are spread across the market, revealing that most listings fall between $500,000 and $1.5 million with a small number of high-value properties extending the upper range. A boxplot with a logarithmic scale was used to compare prices across property types, since the log scale makes it easier to read both affordable and luxury properties on the same chart. A bar chart was also used to show the number of listings per property type, as the boxplot does not show how common each type is. Together, these charts highlight that condos are the most common property type while single family homes have the highest prices.

Property Configuration

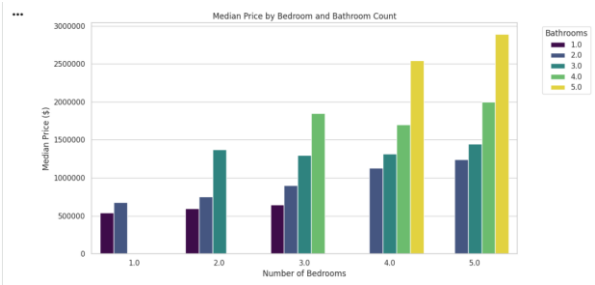

This analysis looked at common inventory layouts and their relationship to price. A grouped bar chart was created showing median price by bedroom and bathroom count, filtered to configurations between one and five bedrooms and bathrooms with a minimum of twenty listings per combination to ensure reliability.

The results show that bathroom count has a stronger association with price than bedroom count. For example, 3-bedroom homes with 4 bathrooms outprice 4-bedroom homes with the same number of bathrooms, suggesting that bathroom count may reflect overall property quality and design rather than size alone.

Feature Drivers

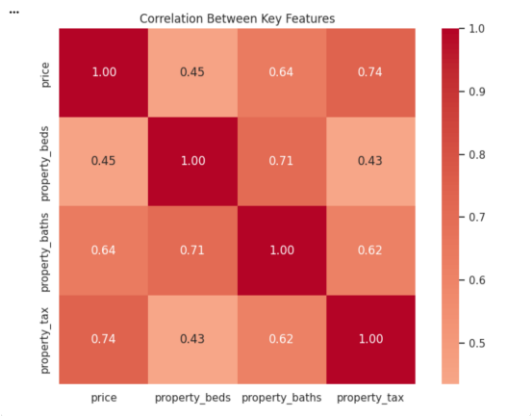

Feature drivers' analysis focused on the relationships between price and key variables including bedrooms, bathrooms, and property tax. A correlation heatmap was created to show the strength of these relationships, and a scatterplot was used to confirm the findings visually. Property tax values were filtered to a realistic range of $500 to $30,000 before the analysis, as the raw data contained corrupt entries with extreme values reaching up to $888 million. It is also important to note that approximately 20% of property tax was estimated during the data cleaning stage using locality medians. As a result, property tax is treated as a supporting indicator in this analysis rather than a definitive price driver.

3.4 Dashboard Implementation

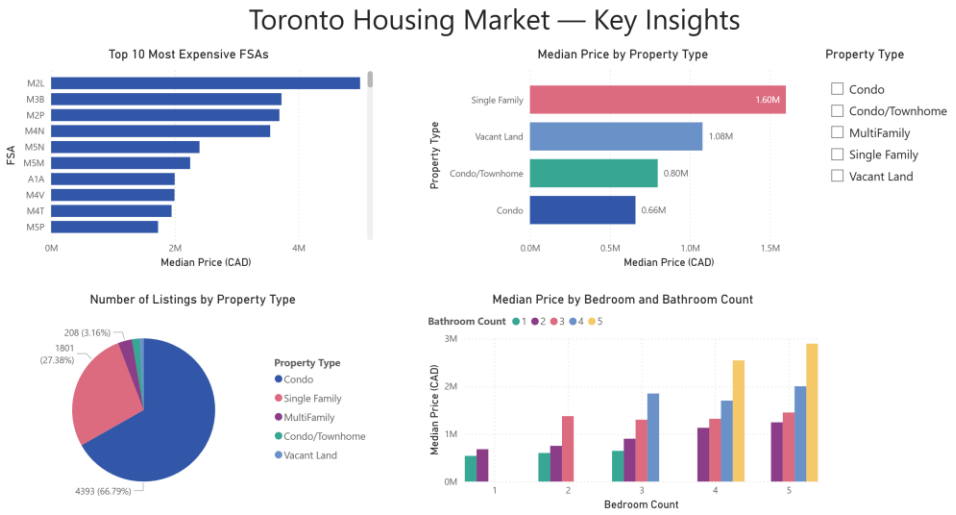

All findings from the analysis were brought together in an interactive Power BI dashboard. The dashboard consists of five pages, including a summary page and four detailed analysis pages covering location, market composition, property configuration, and feature drivers. Users can filter the dashboard by property type, price range, and neighborhood using dynamic slicers, making the analysis accessible and explorable for non-technical stakeholders.

3.5 Process Flow Diagram

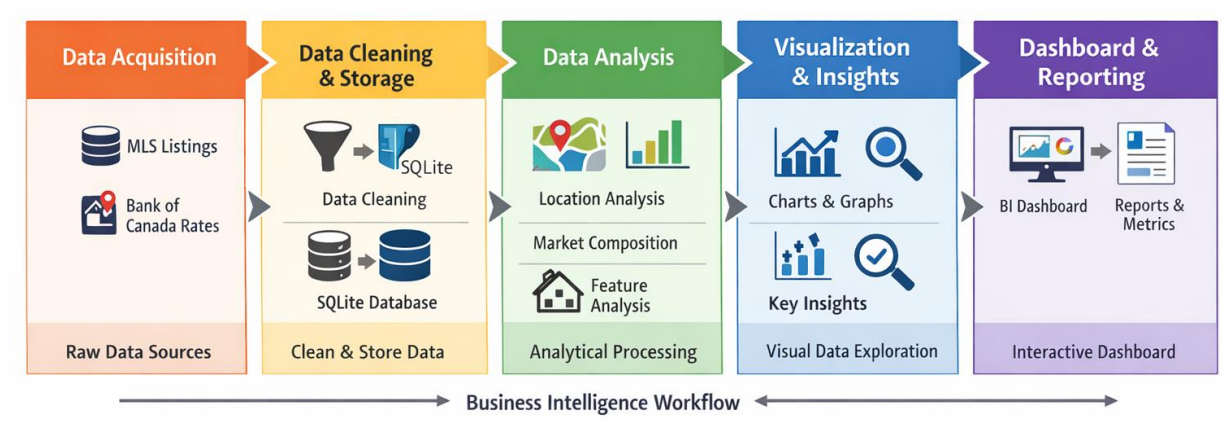

4. Benefits

The Business Intelligence solution developed for the Toronto real estate market delivers substantial, measurable benefits across multiple stakeholder groups, addressing longstanding inefficiencies in how housing data is accessed, interpreted, and acted upon.

Enhanced Decision-Making through Data Consolidation. Prior to this solution, stakeholders in Toronto's housing market were forced to rely on fragmented information sources, manual comparisons, and subjective professional opinions. The BI pipeline—spanning data acquisition from Kaggle and the Bank of Canada API, cleaning via Python, structured storage in SQLite, and interactive visualization through Power BI—consolidates thousands of property listings into a coherent analytical framework. Buyers, sellers, and real estate professionals can now identify pricing patterns by Forward Sortation Area (FSA), property type, and feature configuration within minutes, replacing hours of manual research and reducing the risk of costly misjudgments in a market where median home prices exceed CAD $800,000.

Improved Market Transparency and Equity. One of the defining advantages of this BI implementation is its role in democratizing market intelligence. By surfacing median prices across FSA zones, property types, and bedroom-bathroom configurations in an interactive Power BI dashboard, the solution dismantles the information asymmetry that has historically disadvantaged first-time buyers and individual sellers relative to institutional investors and seasoned agents. The use of median rather than mean pricing further ensures that the insights are not distorted by luxury outliers, producing a more representative and trustworthy picture of market conditions for all participants.

Data-Driven Property Valuation. The diagnostic analysis embedded in this project reveals insights that challenge common assumptions about property valuation. The finding that bathroom count is a stronger price predictor than bedroom count suggests that property quality and spatial design carry more weight than raw size—a nuance that traditional valuation methods often overlook. These evidence-based insights equip real estate agents and sellers with defensible pricing rationale, reducing negotiation friction and improving pricing accuracy. Property tax data, when properly filtered to remove erroneous entries, further strengthens the valuation model as a supporting diagnostic indicator.

Operational Efficiency and Scalability. The structured, five-stage BI pipeline—acquisition, cleaning, storage, analysis, and visualization—establishes a repeatable and scalable framework that reduces manual workload significantly. Tasks that previously required hours of spreadsheet analysis and cross-referencing can be accomplished in minutes through the dashboard's dynamic slicers. Beyond time savings, the SQLite database layer ensures that cleaned data is consistently organized and retrievable, positioning the solution for integration with live MLS data feeds in

future iterations. The modular design of the pipeline also means the framework can be adapted for other Canadian real estate markets, such as Vancouver or Calgary, with minimal rearchitecting.

Macroeconomic Contextualization. The integration of Bank of Canada overnight interest rate data introduces a macroeconomic dimension that purely listing-based analyses cannot provide. Although temporal misalignment between the two datasets precluded a direct statistical merge in this iteration, the contextual overlay of rate trends alongside listing price distributions enables more sophisticated interpretation of market cycles. This benefit is particularly valuable for long-term investors and developers who must assess not only current pricing levels but also the directional influence of monetary policy on housing demand and affordability across Toronto's diverse property segments.

5. Challenges

While the BI solution delivers meaningful value, its development encountered several technical, methodological, and operational challenges that must be acknowledged to accurately assess the solution's current scope and future development requirements.

Data Quality and Integrity Issues. The primary dataset sourced from Kaggle exhibited significant data quality issues that required extensive preprocessing before analysis could proceed. Property tax records contained corrupt entries reaching as high as CAD $888 million—values that clearly reflect data entry errors rather than actual assessments—necessitating threshold-based filtering to a realistic range of $500 to $30,000. Approximately 20% of property tax values required imputation using locality-level medians, introducing measurement uncertainty that limits the variable's reliability as a primary price predictor. Missing values in square footage fields were similarly imputed using median substitution, which, while statistically appropriate, may obscure genuine variance within certain property segments. These preprocessing decisions reflect inherent trade-offs between data completeness and analytical validity that are common in real-world BI deployments.

Geographic Data Inconsistency. The dataset's address locality field contained inconsistent, non-standardized neighborhood labels that could not reliably serve as geographic grouping variables. As a mitigation strategy, Forward Sortation Area (FSA) codes—derived from the first three characters of postal codes—were used as proxy geographic identifiers. While FSA-based grouping provides a more consistent spatial unit, it introduces a granularity limitation: FSA boundaries do not align with recognized Toronto community boundaries such as Rosedale, Leslieville, or Forest Hill. This mismatch reduces the interpretive value of the location analysis for end users who navigate the market using neighborhood names rather than postal code prefixes, representing a user experience gap that future iterations should address through geocoding enrichment or boundary mapping services.

Cross-Dataset Integration and Temporal Misalignment. A significant methodological limitation arose from the absence of standardized listing date fields in the Kaggle housing dataset, which prevented a direct temporal join with the Bank of Canada's Valet API interest rate series. This temporal misalignment confined interest rate data to the role of macroeconomic background context rather than an integrated analytical variable. Consequently, the BI system cannot currently quantify the directional relationship between rate changes and listing price movements—a critical analytical capability for investment forecasting. Resolving this limitation would require either acquiring a timestamped MLS dataset or engineering a date proxy from ancillary fields, both of which involve additional data engineering complexity.

Static Dataset and Lack of Real-Time Refresh. The current implementation operates on a static dataset representing a single historical snapshot of the Toronto housing market. In a market characterized by rapid price fluctuations driven by interest rate changes, immigration patterns, and regulatory interventions, insights derived from a frozen dataset risk becoming obsolete within weeks of publication. A production-grade BI deployment would require an automated Extract, Transform, Load (ETL) pipeline capable of ingesting live MLS data feeds on a scheduled or near-real-time basis, integrated with a version-controlled data warehouse rather than a static SQLite database. The absence of this infrastructure is the most significant barrier to transitioning the current prototype into an operational decision-support system.

User Accessibility and Analytical Literacy. Although the Power BI dashboard is designed to minimize technical barriers, concepts such as median pricing, logarithmic scales in box plots, and correlation heatmaps may still pose interpretive challenges for users without a quantitative background. This analytical literacy gap is a known challenge in BI system adoption, particularly in consumer-facing deployments. Ensuring that the tool is genuinely accessible to first-time homebuyers and independent sellers requires investment in guided tooltips, contextual help text, and user onboarding documentation—components that fall outside the scope of the current prototype but are essential for real-world deployment.

6. Conclusion and Recommendations

This project demonstrates that Business Intelligence methodologies can fundamentally transform how complex real estate markets are understood and navigated. By designing and implementing a structured five-stage BI pipeline—encompassing data acquisition, cleaning, storage, descriptive and diagnostic analysis, and interactive visualization—the team successfully converted a fragmented dataset of Toronto housing listings into a coherent, stakeholder-accessible decision-support system. The resulting Power BI dashboard addresses the core problem of pricing opacity and information asymmetry that has long disadvantaged buyers, sellers, and smaller agencies in Canada's most competitive housing market.

The analytical findings confirm several important market dynamics. Geographic location, measured through FSA codes, remains the most significant determinant of listing price, with substantial variation across Toronto's postal zones. Condominium units dominate the inventory by volume, but single-family detached homes command the highest median prices. Most notably, bathroom count proved to be a stronger price predictor than bedroom count—a counterintuitive finding that underscores the importance of property quality and spatial configuration over raw size in Toronto's premium housing environment. Together, these insights validate the core value proposition of the BI system: that data-driven analysis can surface patterns that subjective, experience-based assessments routinely miss.

Recommendations. To evolve this prototype into a production-grade BI system, four strategic enhancements are recommended. First, the static dataset should be replaced by an automated ETL pipeline that ingests live or regularly refreshed MLS data, ensuring the dashboard reflects current market conditions and supports timely decision-making. Second, the data model should be enriched with standardized listing date fields to enable temporal analysis and allow direct correlation between Bank of Canada interest rate movements and housing price trends—a capability that would substantially increase the system's value for investors and policy analysts. Third, a predictive analytics layer should be developed using regression or machine learning models to generate estimated property valuations based on location, type, bedroom and bathroom configuration, and square footage; this would transform the dashboard from a descriptive reporting tool into a forward-looking decision-support engine. Fourth, a structured pilot deployment should be conducted in partnership with a real estate brokerage to validate the dashboard in operational conditions, gather user feedback on accessibility and interpretation, and refine the interface design prior to broader release. Taken together, these recommendations would position the solution as a scalable, commercially viable BI platform capable of bringing data-driven transparency to stakeholders across Toronto's CAD $90 billion housing market and, with appropriate adaptation, to other major Canadian metropolitan markets as well.

7.References

Yuliia Bulana. (n.d.). *Canada housing dataset*. Kaggle. https://www.kaggle.com/datasets/yuliiabulana/canada-housing

Bank of Canada. (n.d.). *Valet API: Interest rate data*. https://www.bankofcanada.ca/valet/docs

Bank of Canada. (n.d.). *Key interest rate explained*. https://www.bankofcanada.ca/core-functions/monetary-policy/key-interest-rate/

IBM. (n.d.). *What is business intelligence?* https://www.ibm.com/topics/business-intelligence